

& Mathematics

ALEX DAVIES DEEPMIND · UNIVERSITY OF OXFORD

MARC LACKENBY UNIVERSITY OF OXFORD


GEORDIE WILLIAMSON UNIVERSITY OF SYDNEY

$$\lambda_1 \geq \lambda_2 \geq \dots$$

$$\theta_{t+1} = \theta_t - \eta \nabla \mathcal{L}$$

$$p_{\theta}(y | x)$$

ing, every retained head and channel becomes a coordinate in a large discrete allocation space.



04 — **Evidence** How to report speed, MACs, and accuracy cleanly

05 — **Takeaway** What the method actually buys us

BUDGETED SEARCH SPACE, NOT A SINGLE PRUNING RATIO




rank units



move through allocations


...and the government can use it to push the population toward better regions.



Illustrative allocation cells - not a measured embedding

The slide style here keeps dense algorithm text out of the center. Use the upper title for the claim, the middle for structure, and the bottom line for the one sentence people should remember.


sult pages quiet: blurred motifs, compact headers,
readable numbers, no oversized slogans.



Runs

small side note rather than a hidden appendix.

Design rule: one row per experimental decision, no paragraph longer than two lines.




Red curve: reference heuristic.

Fine-tune mismatch The searched mask and physically pruned model do not recover equally.

This layout is for honest engineering slides: name the failure, state its symptom, then explain the fix on the next slide.

oration. It is a consistent system: crisp section breaks, quiet content pages, and mathematical motifs that stay in the background.



contact · repository · paper draft

$$\lambda_1 \geq \lambda_2 \geq \dots$$

$$\theta_{t+1} = \theta_t - \eta \nabla \mathcal{L}$$

$$p_{\theta}(y \mid x)$$